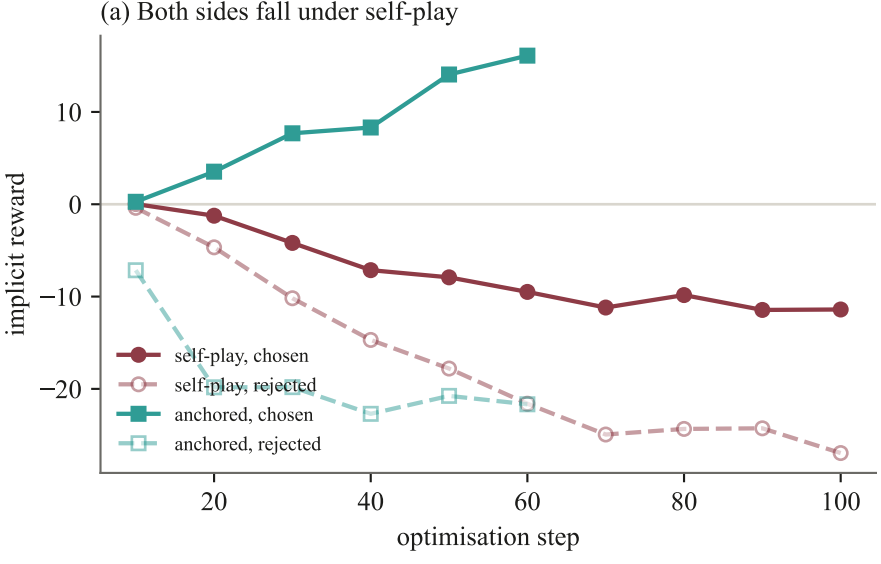
(a) Both sides fall under self-play
implicit reward
10
0
−10
−20
20
40
60
80
100
optimisation step
self-play, chosen
self-play, rejected
anchored, chosen
anchored, rejected

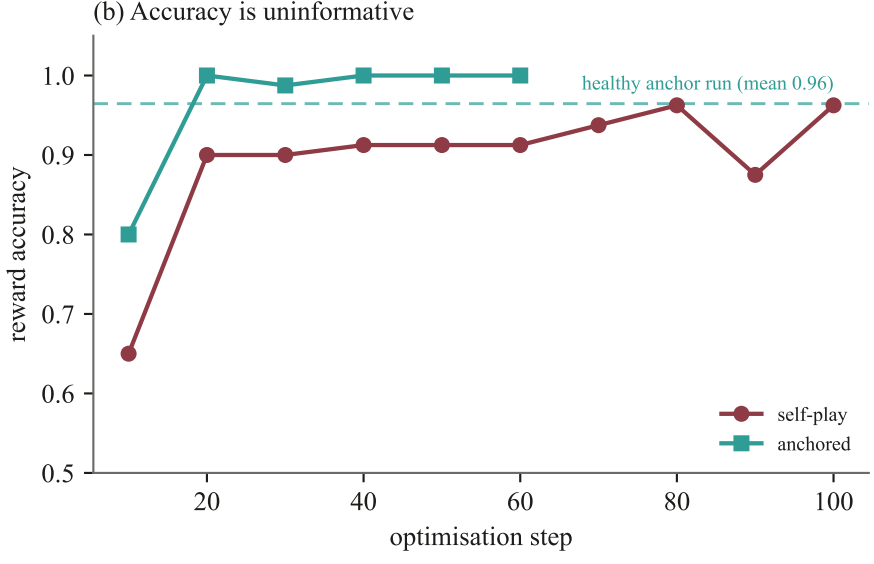
(b) Accuracy is uninformative
reward accuracy
1.0
0.9
0.8
0.7
0.6
0.5
healthy anchor run (mean 0.96)
20
40
60
80
100
optimisation step
self-play
anchored